

```
In [1]: %matplotlib inline
import torch
from d2l import torch as d2l
```

```
In [3]: class LinearRegressionScratch(d2l.Module): #@save
        """The linear regression model implemented from scratch."""
        def __init__(self, num_inputs, lr, sigma=0.01):
            super().__init__()
            self.save_hyperparameters()
            self.w = torch.normal(0, sigma, (num_inputs, 1), requires_grad=True)
            self.b = torch.zeros(1, requires_grad=True)
```

```
In [4]: @d2l.add_to_class(LinearRegressionScratch) #@save
        def forward(self, X):
            return torch.matmul(X, self.w) + self.b
```

```
In [5]: @d2l.add_to_class(LinearRegressionScratch) #@save
        def loss(self, y_hat, y):
            l = (y_hat - y) ** 2 / 2
            return l.mean()
```

```
In [6]: class SGD(d2l.HyperParameters): #@save
        """Minibatch SGD."""
        def __init__(self, params, lr):
            self.save_hyperparameters()

        def step(self):
            for param in self.params:
                param -= self.lr * param.grad


        def zero_grad(self):
            for param in self.params:
                if param.grad is not None:
                    param.grad.zero_()
```

```
In [7]: @d2l.add_to_class(LinearRegressionScratch) #@save
        def configure_optimizers(self):
            return SGD([self.w, self.b], self.lr)
```

```
In [8]: @d2l.add_to_class(d2l.Trainer) #@save
        def prepare_batch(self, batch):
            return batch

@d2l.add_to_class(d2l.Trainer) #@save
def fit_epoch(self):
    self.model.train()
    for batch in self.train_dataloader:
        loss = self.model.training_step(self.prepare_batch(batch))
        self.optim.zero_grad()
        with torch.no_grad():
            loss.backward()
            if self.gradient_clip_val > 0: #나중
                self.clip_gradients(self.gradient_clip_val, self.model)
            self.optim.step()
        self.train_batch_idx += 1
    if self.val_dataloader is None:
        return
    self.model.eval()
    for batch in self.val_dataloader:
        with torch.no_grad():
            self.model.validation_step(self.prepare_batch(batch))
        self.val_batch_idx += 1
```

```
In [9]: model = LinearRegressionScratch(2, lr=0.03)
data = d2l.SyntheticRegressionData(w=torch.tensor([2, -3.4]), b=4.2)
trainer = d2l.Trainer(max_epochs=3)
trainer.fit(model, data)
```



```
In [10]: with torch.no_grad():
          print(f'error in estimating w: {data.w - model.w.reshape(data.w.shape)}')
          print(f'error in estimating b: {data.b - model.b}')
```

```
error in estimating w: tensor([ 0.1526, -0.2056])
error in estimating b: tensor([0.2406])
```

Discussion: Before diving into the main content, I had maintained the stance of not typing out comments, which is why I didn't include `#@save`. However, since it kept appearing repeatedly, I decided to look it up and learned that it indicates sections of code that can be reused multiple times. Thus, unlike other comments, I decided to include it. This section focused on implementing a linear regression model and training the model's weights and bias using SGD. I'm curious how the learning rate (`lr`) in the SGD class will affect the final result, as this hasn't been covered yet. The loss function used is MSE, which I was happy to see because it's a function I learned about in computational mathematics before. During the training process, I came across the new concept of an "epoch." After the data was trained, the estimated parameters and the actual parameters were compared by printing out the error between them.

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js